

Data Engineering Lifecycle: It describes how data is captured, stored, processed, and delivered. Main stages: Generation → Ingestion → Storage → Transformation → Serving. Undercurrents (always present): Security, Data Management, DataOps, Data Architecture, Orchestration, Software Engineering.

Type A Data Engineers (Abstraction): They keep architecture simple and well-managed services and off-the-shelf tools. Their focus is running the lifecycle efficiently without building custom systems. **Type B Data Engineers (Build):** They build custom data systems for scale or unique, mission-critical needs when existing services cannot meet the requirements.

Internal-Facing Data Engineer: They focus on building and maintaining data systems that support internal teams such as analytics, BI, and machine learning. Their work improves data access, quality, and reliability inside the organisation. **External-Facing Data Engineer:** They build data systems that power products used by customers or partners. This includes handling application data, APIs, and features where data engineering directly affects external users.

Upstream: Software Engineers, Data Architects, DevOps. **Downstream:** Data Analysts, Data Scientists, ML Engineers, Business Leaders (CIO/CTO/CDO).

Data Pipeline: A system that collects, processes, and moves data from sources to storage or applications. It often cleans, transforms, and organizes data so it can be effectively used for analytics, reporting, or machine learning. **Key Stages:** Data Sources → Ingestion → Storage → Transformation → Serving → Analytics/ML.

Storage Considerations: No one-size-fits-all. Different data types and workloads need different storage solutions. Schema flexibility: Some systems need fixed schemas; others handle evolving or semi-structured data. Storage vs. query: Some store data cheaply, others let you query it directly. Performance and scalability: Consider read/write speed, ability to scale, and reliability guarantees (SLAs).

Data acquisition/question: process of collecting raw data from various sources and bringing it into a system for storage and processing. It's often the main bottleneck in a data pipeline due to volume, velocity, and diversity of data. **Considerations:** Data is often collected for multiple purposes? Batch vs. Streaming and Method: Push or pull? On-demand ingestion needed. Volume and Velocity: How much data and how fast does it arrive. Data Format: What format is the incoming data in?

Data cleaning and transformation: process of converting raw data into a usable, high-quality form for downstream applications. This stage is where the data starts generating real value for users or systems. **Key Considerations:** Data Quality: Does the data meet downstream quality requirements? System Compatibility: Can existing systems handle the current format, or are transformations needed? Simplicity: Are transformations simple and self-contained? Cost vs. ROI: What is the cost and return on investment of the transformation?

File System Navigation: pwd → current directory, cd name → change directory, cd .. → previous directory, ls [name..] → list contents, ls -al → detailed list, ls -R → recursive list, mkdir [name..] → create directory, mv [old..] .. → move/file, cp [source..] [destination..] → copy file

File Content: can filename → show file content, most recent file_name → page-by-page view, head filename → first 10 lines, tail filename → last 10 lines, w filename → lines, words, characters, sort filename → sorted content

Piping: Treats everything as a stream of data. Output of one command can be passed directly to another using —. Example: cat file_name → sort → head -3 → reads, sorts and shows the first 3 lines.

Grep: A command-line tool used to search for patterns in files and display matching lines. **LWK:** A text-processing language for extracting, transforming, and manipulating data from files line by line.

Data Acquisition: File Access → Data from existing files or downloaded datasets. Formats: CSV, Excel, XML, JSON, text, images. **Database / Application Access:** Data fetched programmatically via APIs or scraping. Formats: XML, JSON. **Change Data Capture (CDC) Logs:** Data from messages, streams, or publish/subscribe systems. Formats: JSON, log files, event streams. **File-Based Approach:** Data stored in files for local analysis. Formats: CSV, spreadsheets; analyzed via formulas, charts, or Python/pandas.

Push: Source actively sends data to target as it happens. Difference: Source initiates transfer. Ex: Notifications, Kafka, MQTT. **Pull:** Target requests data from source when needed. Difference: Target initiates transfer. Ex: File download, database query, API call. **Tell:** Target repeatedly checks source for updates at intervals. Difference: Target checks periodically rather than continuous or on-demand. Ex: Web crawling, periodic database checks.

Bounded Data: Fixed-size, finite datasets processed as a batch. Ex: CSV files, API data pulls. **Unbounded Data:** Continuous, infinite data streams arriving over time; must be windowed or batched for processing. Ex: Sensor readings, log streams.

Latency: Time delay between data generation and its availability for user; critical for real-time streams. Ex: Time from a sensor reading to dashboard update. **Scalability:** System's ability to handle increasing data volumes or bursts without performance loss. Ex: Adding nodes to a cluster to manage higher traffic.

ETL (Extract → Transform → Load): Extract data from source (batch-pull or push), Transform/clean before loading, Load into storage (e.g., data warehouse)

ELT (Extract → Load → Transform): Load raw data into a lake or warehouse, Transform/clean for each use case within the storage layer

ETL Process: Extract: Data collected from multiple sources. Methods: Custom scripts using APIs, open-source or SaaS products. Examples: Ad platforms (Facebook Ads, Google Ads), backend databases, sales CRMs. Challenges: Requires technical skill, maintenance of extraction scripts. **Transform:** Purpose: Normalize and model raw data. Methods: Programming languages (Python, Scala), technologies (Spark, Hadoop), or GUI-based ETL platforms. Considerations: Remodeling data to fit specific requirements like a data warehouse for BI and end-user access. **Load:** Purpose: Load transformed data into target system. Data is not copied to the warehouse.

ETL Pros: Ensures data is cleaned and transformed before entering the warehouse, maintaining quality. Supports complex batch logic and transformations outside the warehouse. Reduces storage of unnecessary raw data. **Cons:** Requires separate transformation infrastructure, adding complexity. Slower for very large datasets due to pre-loading transformations. Less flexible for ad-hoc queries or changing analytics needs.

ELT Pros: Leverages warehouse processing power, scalable for large datasets. Preserves raw data, allowing reprocessing or different analyses. Faster initial loading since transformations happen after loading. **Cons:** Warehouse must be powerful enough to handle transformations. Raw data storage can be large before transformation. Managing transformations inside the warehouse can be complex.

Undercurrents in Data Acquisition: Security: Data ingestion can expose pipelines to risks like unauthorized access or man-in-the-middle attacks. Use VPNs, private connections, or encryption in transit. Data Ethics, Privacy, and Compliance: Only collect necessary data, especially sensitive/personal information. Prefer anonymized or masked data for testing and training. **Data Management:** Handle schema changes and ensure you are notified when source data structure changes. **DataOps:** Methodology to streamline and automate data pipelines, ensuring reliable, fast, and collaborative data delivery. **Data Architecture:** Blueprint for storing, organizing, and managing data to enable consistency, scalability, and accessibility throughout its lifecycle. **Orchestration:** Schedule and manage complex data pipeline tasks as complete workflows, not individual jobs. **Software Engineering:** Use code reviews, test coverage, and automated tests to simplify and secure pipelines.

Data Quality: Data quality depends on users; errors are common. Metadata often separate, hard to keep in sync. Backup and sharing are manual; multiple copies create redundancy. Hard to prevent inconsistent changes or maintain data integrity.

Data Quality Attributes: Believable: Trusted, credible, and regarded as true. Value-added: Provides benefits or advantages when used. Relevant: Useful and applicable for the task. Accurate: Correct, reliable, and error-free. Interpretable: Clear, understandable, and well-defined.

Data Cleaning: **Missing Values:** Cause: Not recorded (survey non-response, sensor faults). Impact: Can bias results if not random. Detection: Use built-in functions (e.g., Pandas). Remedy: Remove or impute carefully.

Default Values: Cause: Set values instead of blanks (e.g., "NA", "-1"). Impact: Can distort analysis (e.g., averaging text values). Detection: Domain knowledge, metadata, frequency checks. Remedy: Treat as missing values.

Incorrect Values: Cause: Device errors, unreliable responses, data entry mistakes. Impact: Distorts accuracy. Detection: Domain knowledge, statistical tests (e.g., Benford's law). Remedy: Remove or correct if possible.

Inconsistent Values: Cause: Free-text input variations (e.g., "Röhm" vs "Roehm"). Impact: Underrepresentation or miscalcification. Detection: Cluster analysis or variant detection. Remedy: Reconcile to a single standard value.

Handling CSV in Python: The csv module has csv.reader which returns rows as lists and csv.DictReader which returns rows as dictionaries. Manual type conversion may be needed. Pandas provides pd.read_csv which reads data into a DataFrame, handles missing data, data types, filtering, and basic statistics, and is preferred for analysis.

Python Missing Values: Use the na_values parameter in pd.read_csv() to mark placeholders. After import, dropna() removes rows with missing values, fillna() fills missing values, and replace() changes specific values. **Type Conversion:** Use int(), float(), and datetime.strptime() in Python. In Pandas, astype() converts column types. **Data Visualization:** Use Matplotlib for bar, line, and scatter plots. Pandas also allows quick plotting directly from DataFrames or Series.

DBMS: is software that stores, manages, and allows efficient retrieval, insertion, and updating of data in a structured way. Pros: Centralized management ensures data security, integrity, and consistency. Supports multiple users and concurrent access efficiently. Cons: Can be resource-intensive and complex. Administer. May require specialized knowledge to design and maintain.

Relational Database (Structured Data): A type of DBMS that stores data in relations (tables) with rows (tuples) and columns (attributes). Each table has a schema defining its structure and data types, and tables can be related via keys. Pros: Data is stored in a normalized form, supports joins through constraints and ACID transactions. Powerful querying with SQL supporting complex joins and aggregations. Cons: Rigid schema; changing structure requires altering tables. Scaling horizontally is difficult; large joins can impact performance.

Primary Key: A unique attribute or combination of attributes that identifies each row in a table. Cannot be NULL. **Foreign Key:** An attribute in one table that refers to the primary key in another table, creating a link between the two tables.

Schema Diagrams: Graphical representations of a relational database showing tables, columns, and relationships. Normalised schema avoid data redundancy by storing each fact only once. Includes notations like Entity-Relationship Diagrams (ERDs).

Schema Design: Automatic schema generation: Creates tables from DataFrames or CSVs but usually lacks foreign key constraints, and proper normalization. Manual schema design: Define tables with SQL DDL (CREATE TABLE) to control data types, relationships, and integrity; ensures consistency but makes loading raw data harder. Practical approach: Use a normalized schema with code-based loading to balance control and ease of use.

Relational vs Non-Relational: Relational uses fixed schemas; NoSQL allows flexible or dynamic schemas. Relational stores structured data in tables: NoSQL handles unstructured or semi-structured data. Relational enforces ACID transactions; NoSQL does not.

Analytical vs Operational: Analytical databases are optimized for complex queries and reporting; operational databases handle real-time transactions. Analytical stores historical or aggregated data; operational stores current transactional data. Analytical focuses on read-heavy workloads; operational focuses on frequent inserts, updates, and deletes.

Commercial vs Open-Source: Commercial databases require licenses and paid support; open-source databases are free and community-supported. Commercial offers proprietary features and enterprise optimizations; open-source relies on community contributions. Commercial provides enterprise-grade security and SLAs; open-source may need extra setup for similar guarantees.

Disk-Based vs Main-Memory vs Cloud: Disk-based stores data on disks; main-memory stores data in RAM for faster access; cloud-based is hosted remotely and scalable. Disk/main-memory rely on local infrastructure; cloud provides managed services and high availability. Disk-based is slower for large queries; main-memory is faster but limited by RAM; cloud offers flexible scaling.

Quering PostgreSQL from Python: Use a connection (conn) to execute SQL statements. Pandas can load query results into a DataFrame using pd.read_sql(conn). Use a utility function, for example query(), to distinguish between tables and queries that return data.

Building Data from Storage: Load an entire table from the database into Python for processing, for example using pd.read_sql(table_name, 'measurements', conn). **Pros:** Easy to use, works across systems, and gives full control in Python. **Cons:** Cannot leverage database optimizations and must fit the data in memory. Advanced database features, such as user-defined functions or stored procedures, can improve performance but may cause platform lock-ins.

Data Storage Approach 1: External GUI or command line tools: Load data directly from files using tools like psql or SQL browsers such as pgAdmin. **Pros:** Fast, straightforward, no programming needed. **Cons:** Only one-to-one CSV-to-table mapping, stops at the first error, no data transformation. **Approach 2: Programmatically using SQL:** Generate SQL INSERT statements via scripts or code, for example Python, to load data. **Pros:** Flexible, allows data cleaning and transformations, can be optimized for the use case. **Cons:** Requires development time; schema or transformation changes must be handled manually. **Approach 3: Pandas loading code:** Use a Pandas DataFrame and the to_sql() function to load data into a database. **Pros:** Flexible, integrates Pandas data cleaning and transformations. **Cons:** Needs custom coding, set of tables and constraints may be required.

Operational Systems (OLTP): Purpose: Handles day-to-day transactions. Data Volume: Small, current data. Data Type: Simple, short queries, Updates: Frequent inserts, updates, deletes. Design: Highly normalized. Example: Banking system.

Analytical Systems (OLAP): Purpose: Supports analysis and decision-making. Data Volume: Large historical data. Data Type: Complex, long queries, Updates: Mostly read-only. Design: Denormalized for faster querying. Example: Retail sales reporting: Extract sales data for the past year, transform, group by product category, Aggregate: calculate total revenue per category, Sort results by revenue, Load into reporting dashboard.

OLAP Operations: Roll-up: Summarize data along hierarchy (e.g., city → country). Drill-down: View more detailed data (e.g., year → month). Slice: Select a single dimension (e.g., products × cities for January).

Dice: Select multiple dimensions to create a smaller cube. Pivot: Rotate dimensions for a new perspective (e.g., switch rows/columns).

Data Warehouse: A separate database for analytics that combines data from multiple operational systems. It stores data in a native form until needed for analysis. Pros: Can handle large volumes and diverse types of data. Flexible and supports complex queries. It stores raw, structured, semi-structured, and unstructured data. Structured and unstructured data. **Cons:** Data is denormalized and lacks referential integrity. It supports joins with descriptive attributes. **Snowflake Schema:** A schema where dimension tables are normalized into multiple related tables around the central fact table. Reduces redundancy but makes queries slightly more complex.

Data Lake: A centralized repository that stores raw, structured, semi-structured, and unstructured data in its native form until needed for analysis. Pros: Can handle large volumes and diverse types of data. Flexible and supports complex queries. It stores raw, structured, semi-structured, and unstructured data. Structured and unstructured data. **Cons:** Data is denormalized and lacks referential integrity. It supports joins with descriptive attributes. **Snowflake Schema:** A schema where dimension tables are normalized into multiple related tables around the central fact table. Reduces redundancy but makes queries slightly more complex.

Data Lake: A centralized repository that stores raw, structured, semi-structured, and unstructured data in its native form until needed for analysis. Pros: Can handle large volumes and diverse types of data. Flexible and supports complex queries. It stores raw, structured, semi-structured, and unstructured data. Structured and unstructured data. **Cons:** Data is denormalized and lacks referential integrity. It supports joins with descriptive attributes. **Snowflake Schema:** A schema where dimension tables are normalized into multiple related tables around the central fact table. Reduces redundancy but makes queries slightly more complex.

Data Warehouse: A structured repository that stores raw, structured, semi-structured, and unstructured data. It supports joins with descriptive attributes. Pros: It stores raw, structured, semi-structured, and unstructured data. Structured and unstructured data. **Cons:** Data is denormalized and lacks referential integrity. It supports joins with descriptive attributes. **Snowflake Schema:** A schema where dimension tables are normalized into multiple related tables around the central fact table. Reduces redundancy but makes queries slightly more complex.

Data Lake: Stores raw, structured, semi-structured, and unstructured data. Schema-on-read, flexible for multiple analytics and machine learning use cases. **Cons:** Harder to enforce data quality, governance, and security. Query performance can be slower than structured warehouses.

Data Warehouse: Stores raw, structured, cleaned, and processed data. Optimized for analytics and machine learning use cases. Can handle large volumes but may have slower query performance and harder governance.

Web API (Application Programming Interface): is a set of rules and protocols that allows applications to communicate over the internet. It enables one software system to request data or services from another using standard web protocols like HTTP.

Web Scraping: The process of extracting data from websites, either by downloading files or parsing data embedded in web pages. **General Approach: Reconnaissance:** Identify the source, structure, and terms of service. **Webpage Retrieval:** Download pages, often using scripts that generate URLs. **Data Extraction:** Parse content and extract raw data. **Cleaning and Transformation:** Convert data into the required format. **Storage and Processing:** Store data in a structured format for further analysis. **Analysis:** Save or combine with other datasets. **Notes:** Web pages are written in HTML, which may be semi-structured or poorly formatted. Browsers are forgiving with HTML, but parsers may fail on incorrect markup. **Tools and Libraries:** Used commonly (curl, wget, Python's pandas.read_csv(), BeautifulSoup, Scrapy, Import.io, Google Sheets ImportHTML, and commercial web-crawling services).

HTTP Headers: Pay attention to URL headers, tokens; they may be looped over for multiple datasets (e.g., monthly or yearly feeds). Look for patterns in links to automate data access efficiently. Check for access tokens or API keys if required. Inspect the web page structure to locate the desired elements. Use browser page inspector to identify HTML elements for extraction. Complex structures may require tokenizing child nodes, which can be handled with libraries like BeautifulSoup in Python.

Web Page Retrieval in Python: Request Types: GET – fetches a webpage, optionally with parameters. POST – sends data to a web form or endpoint. Python Requests Library: Simple page: request.get(URL). Page with parameters: requests.get(URL, params=your_value). POST form: requests.post(URL, params=value). URLS: Uniform Resource Locator, which is the address of a resource. Format: protocol://site/path/to/resource. Protocol can be http, https, or ftp. Site is the domain or IP (optional port). Path is the resource location and may include query parameters.

Website Crawling: Scrapy: Python framework for multi-page crawling (spiders follow links and extract data). Selenium: Programmable browser to simulate user actions, useful for interactive pages and running JavaScript.

Selecting Content in a Webpage: Text Patterns: Simple string matching; limited for complex structures. DOM Navigation: Traverse the document object model to find elements. CSS Selectors: Use tags, classes, and IDs to select elements; easy but depends on consistent CSS usage. XPath Expressions: Powerful way to navigate the document tree and select nodes or subtrees; supports filters on attributes and values.

CSS Selectors: Class: table, data selects elements with class data.ID: #results or div#results selects the element with ID results. Position: :first-child, :last-child, :nth-child(n, :nth-type(n)) selects elements by order.

Web API Programming: Getting Data via APIs: Many websites provide APIs to request data programmatically. Common formats: JSON and XML.

Web Service Standards: RESTful APIs: Stateless operations using standard HTTP methods (GET, PUT) with responses in HTML, XML, or JSON. XML Web Services (SOAP): XML-based messaging standard, more complex and error-prone than REST.

JSON: Purpose: Lightweight, semi-structured data exchange; widely used in APIs. Structure: Nested key-value pairs, arrays, and objects. Syntax: → objects, [] → arrays; flexible and easy to parse. Example: {"name": "John", "age": 21, "courses": ["Math", "Physics"]}

XML: Purpose: Display content in a browser; focuses on layout and presentation. Tags: Predefined (e.g., **p**, **div**, **img**). Structure: Not strict; mainly for human-readable pages. Example: <p>Welcome to the website!</p>

DTD Document Type Definition (DTD): XML documents let you own your tags, but a DTD defines which tags are allowed. A DTD acts like the grammar of the XML document. It can specify: Which elements exist? Their allowed attributes; their nesting and order; which elements must appear. You can link a DTD to an XML document so the document can be validated against it. Well-formed XML: Follows XML syntax rules (proper nesting, matching tags, one root). Valid XML: Well-formed and follows a DTD/schema.

How to Query or Filter XML: DOM navigation: Walk through the XML tree programmatically using the Document Object Model. XPath: A concise way to select nodes, attributes, or entire subtrees inside one XML document. XQuery: Built-in XPath: a full query language for querying multiple XML documents. CSS Selectors: Mainly for HTML, not general XML.

XPath Document Model (Tree View): XPath treats an XML document as a tree. A special synthetic root node sits above the document's actual root element. Internal nodes = XML elements.

Semi-Structured Data: Data that does not follow a fixed schema but carries tags or keys to describe its structure. It can have nested, hierarchical organization and mixed data types. Pros: flexible schema allows for more complex data.

Relational Structured Data in Databases: Relational DBMS Support – Some SQL databases like PostgreSQL can store and query JSON and XML data via standard XML functions (e.g., XML::store(), XML::extract()). NoSQL Databases – MongoDB is designed for JSON storage and retrieval, handling flexible, nested structures efficiently. EX: JSON: "name": "Bob", "degree": "BSc (CS)", "courses": [...] stored in a JSON/SON column.

NoSQL Types: Document Stores – Nested key-value pairs, e.g., MongoDB, Redis. Pros: Highly scalable; optimized for horizontal scaling across servers. Flexible schema supports evolving or unstructured data. Cons: Weak consistency guarantees (eventual consistency in many systems). Complex queries and joins can be less efficient or require workarounds.

NoSQL Data Types: Document Stores – Nested key-value pairs, e.g., MongoDB, XML databases. Column Stores – Large-scale column-oriented storage, e.g., BigTable, HBase. Key-Value Stores – Simple key-value mapping, e.g., Redis, DynamoDB, Cassandra. Graph Databases – Nodes and relationships, e.g., Neo4j. No standard model or API across NoSQL systems.

Schema-on-Read: Data can be ingested without predefined schema. Structure is inferred when reading/querying. Flexible: good for semi-structured or evolving data. Common: XML, JSON, YAML.

Schema-on-Write: Data must be defined before inserting data. Structure is fixed and enforced by the database. Ensures strong integrity and consistency. Relational databases like PostgreSQL or MySQL.

MongoDB: Flexible schema: documents can differ in structure. Collections correspond to tables, documents to rows, and fields to columns. Pros: Flexible schema; easy to handle evolving or semi-structured data. Supports nested documents and arrays, reducing the need for joins. Cons: Complex relationships between documents can be harder to query efficiently. Weaker consistency guarantees, with eventual consistency in shared setups.

Querying and Operations: Create: insert SQL documents into collections. Read: find() with filters, projections, sorting supports predicates (get(), or, etc.). Flexible Queries: implicit equality and conjunction; can combine conditions. Ex: Document: "name": "Sue", "age": 26, "status": "pending". Query: Find students older than 20 in Sydney.

MongoDB Aggregations: Pipelines of stages for filtering, grouping, sorting, computing. Composes complex operations in one flow. Used for analytics and reporting. Scalability: Auto-sharding; partitioned collections; sharded servers. Single-leader replication for updates; followers read (eventual consistency). Load balancing and high availability via distributed architecture. Ex: Aggregate sales totals per region; student data split across shards.

Graph Database: Stores data as nodes (entities), edges (relationships), and properties (attributes). Properties: Attributes or metadata.

Neo4j (Graph Database): Property graph model: nodes (entities), edges (relationships), and properties (attributes). Flexible schema: nodes and edges can have different properties. Features include ACID compliance, high availability, and the Cypher query language. Advanced features: indexing, full-text search, graph algorithms, and visualization using Neo4j Bloom. Example: {:(Person)-[FACTED_IN]-(Movie)} represents (Person)-[:FACTED_IN]->(Movie).

Spatial Data: Data is interpreted as points in space. Structure: Geographical coordinates (longitude, latitude). Examples: GPS locations, coordinates in a map, coordinates in a grid. Spatial Data: Specifically identifies geographic locations or boundaries on Earth (e.g., cities, suburbs).

OpenStreetMap API: Provides JSON data with latitude, longitude, bounding box, and place details.

SpatialDB (Spatial Database Management System): Stores large amounts of spatial data. Built-in spatial semantics in queries. Specialized indexing for fast spatial access. GIS (Geographic Information System): Client-side DBMS focused on spatial data. Pros: Good for spatial data storage and retrieval.

Spatial Data Types: Point Data: Point: Data represents individual points in space. Example: GPS locations from raw data; pixels in satellite images, feature vectors from text. Region Data: Represents objects with spatial extent (area or boundary). Line Data: Typically stored as vector data using polygons, line segments, etc.

Models of Spatial Information: Object-Based (Entity-Based) Model: Represents distinct, identifiable objects relevant to an application. Objects have attributes (spatial: location; non-spatial: name, type, etc.) and operations (functions on attributes). Example: Road objects have location, type, number of lanes; operation: calculate length or intersection. Field-Based (Space-Based) Model: Represents spatial phenomena as continuous fields over space. Example: Forest stand map viewed as continuous field rather than individual polygons.

WGS84: Standard geographic datum used by the global GPS system. **GD94:** Official geodetic datum for Australia, based on ITRF but fixed to reference points in Australia. **Differences:** Both are geocentric and use almost identical spheroids, so coordinates are mostly interchangeable for mapping and GIS.

PostGIS: PostgreSQL extension for spatial data. Geometry: planar shapes; Geography: spherical, large-scale. Supports spatial functions (area, distance, intersection) and R-Tree indexing. Import/export via GeoJSON, KML, coordinate transformations supported.

Spatial Operations: Set-Based: union, intersection, containment. Topological: touches, disjoint, overlap, directional. Directional: north, east, etc. Metric: distances, lengths.

Topological Relationships: Relationships unchanged under deformation (no tearing/merging). Interior, boundary, exterior define object shape. Quies: e.g., relationship between A and B, or find objects related to A. Nine-Intersection Model: 3x3 matrix of object intersections (inner, outer, boundary). Pros: fast. Cons: 1 = empty, 1 = intersection. Defines 8 possible 2D relationships for objects without holes.

GeoPandas: Python library for geographic data, building on Shapely, Fiona, and GeoJSON. Supports complex geo-features (Shapefiles, MultiPart, MultiLine, MultiPolygon). GeoDataFrame: Table-like structure built on GeoSeries. GeoSeries: Series of geometric objects (Point, Line, Polygon, MultiPoint, MultiLine, MultiPolygon). GeoDataFrame: Table-like structure built on GeoSeries, applies attribute + geometry.

GeoSeries: Attributes include area, bounds, total_bounds, geom_type, and crs. Methods include distance(), centroid_to(crs), and plot(). Relationship tests include contains(), intersects(), and to_crs(). **GeoDataFrame:** Combines tabular data with one special geometry column (GeoSeries). Spatial methods act on the geometry column. Supports attribute joins (merge()) and spatial joins (join()).

Plotting: Use the plot() function for GeoSeries or GeoDataFrame. You can customize colors, markers, legends, and overlay multiple layers (all layers must have the same CRS). **Reading GeoSpatial Data:** Shapefile, MapInfo, and GeoJSON can be read using read_file(). For PostGIS, use read_postgis(). Which requires a database connection.

Geo-aware Web APIs: RESTful APIs allow fast responses for modern and mobile applications. Can return spatial data in JSON or XML formats. Useful for location-based services like maps, navigation, and local search.

Spatial Data Exchange Formats: GEOSON: JSON-based format for geographic data; supports points, lines, polygons, and multi-geometries, along with properties. Example: Representing a location with coordinates and a name. KML (Keyhole Markup Language): XML-based format for maps, placemarks, polygons, and 3D models; used by Google Earth and mapping software. Example: Describing a city's location and details for map visualization.

MapInfo TAB: ESRI's OpenStreetMap Standard for exchanging map data. Other OGC standards: WMS (Web Map Service), WFS (Web Coverage Service). Proprietary formats: MapInfo TAB, ESRI Shapefile.

Temporal Data: Data marked with time, like sensor readings, medical records, or transactions. Helps analyze trends and sequences. Historic (Long-term): Data examining data over long periods (10-30 years) reveals patterns and insights missed by visual analysis.

Reasoning: Data reasoning with time, like temporal joins, and temporal queries. Helps analyze trends and sequences. **Pros:** Enables historical analysis and auditing. Supports reasoning with past, present, and future data. **Cons:** More complex to design and maintain than conventional databases. Can have higher storage and processing overhead.

Temporal Data Types: Instant types: Represent a single point in time. DATE: Day, month, year. TIMESTAMP: Date + time with fractions of a second. TIME: Hours, minutes, seconds (not date). Interval types: Represent a duration of time. Example: Year-Month intervals (e.g., 1 year 2 months). Period Data Type: Represents a time span with a start and end date/time. **Period:** Used for validity or active periods.

Kinds of Time: User-defined time: simple, uninterpreted value (e.g., birthdate). **Valid time:** when a fact is true in reality; can move forward/backward. **Transaction time:** when a fact is stored in the DB; only moves forward, enables read-after-write consistency.

Non-temporal Table: Store and manage data with respect to time, supporting time-based queries. Include transaction-time and valid-time-based queries. **Temporal Table:** Store and manage data with respect to time, storing historical, current, and future information, and supports time-based queries. **Pros:** Enables historical analysis and auditing. Supports reasoning with past, present, and future data. **Cons:** More complex to design and maintain than conventional databases. Can have higher storage and processing overhead.

Temporal Data Types: Instant types: Represent a single point in time. DATE: Day, month, year. TIMESTAMP: Date + time with fractions of a second. TIME: Hours, minutes, seconds (not date). Interval types: Represent a duration of time. Example: Employees (fn, name, city, since) → [since, now) represents current validity.

Valid-Time Tables: History: Track periods of validity using two DATE columns, start_date and end_date, intervals defined by start_time and end_time. Example: InPosition (fn, pos, start_date, end_date), intervals defined by start_time and end_time. **Time-series:** Generalized Time Tree: Store and manage data over time. Examples: Weather data, trip logs, road usage, sleep patterns, cymometry/panel data. Workflow: Ingress: File import, database access, web scraping. Storage: Standard file formats or databases. Analysis and Reporting: In-database or via applications (e.g., Python), plus visualization.

Timelines: Data Storage: Approaches include point-based storage, where each row represents one time point. This is simple, easy to compare and index, and may require multiple rows if value repeat. Sequence-based storage uses a single row to hold arrays of time points, which requires array operations such as UNNEST and array_agg. Dedicated time-based databases include TimescaleDB, InfluxDB, and SciDB.

Point-Based Representation: Stores each timestamp as a separate, atomic row. Supports comparisons, filtering, and joins efficiently. Ideal for OLAP queries and analytics. Uses B-Tree indexing for fast access to individual time points. Working: Query a specific timestamp or range easily, e.g., sum of sales per day.

